

CSE 380: HW 0

Evan Ott

UT EID: eao466

September 30, 2016

0.0

See Figures 1 and 2 along with Tables 1 and 2.

The first set of data could indicate plausibly one of two things: Italy and Saudi Arabia are competitive and want to have the best supercomputers, or that there are (or were) features at the time in the country that make it easy to build the best supercomputer (cheap electricity, increased funding, etc.).

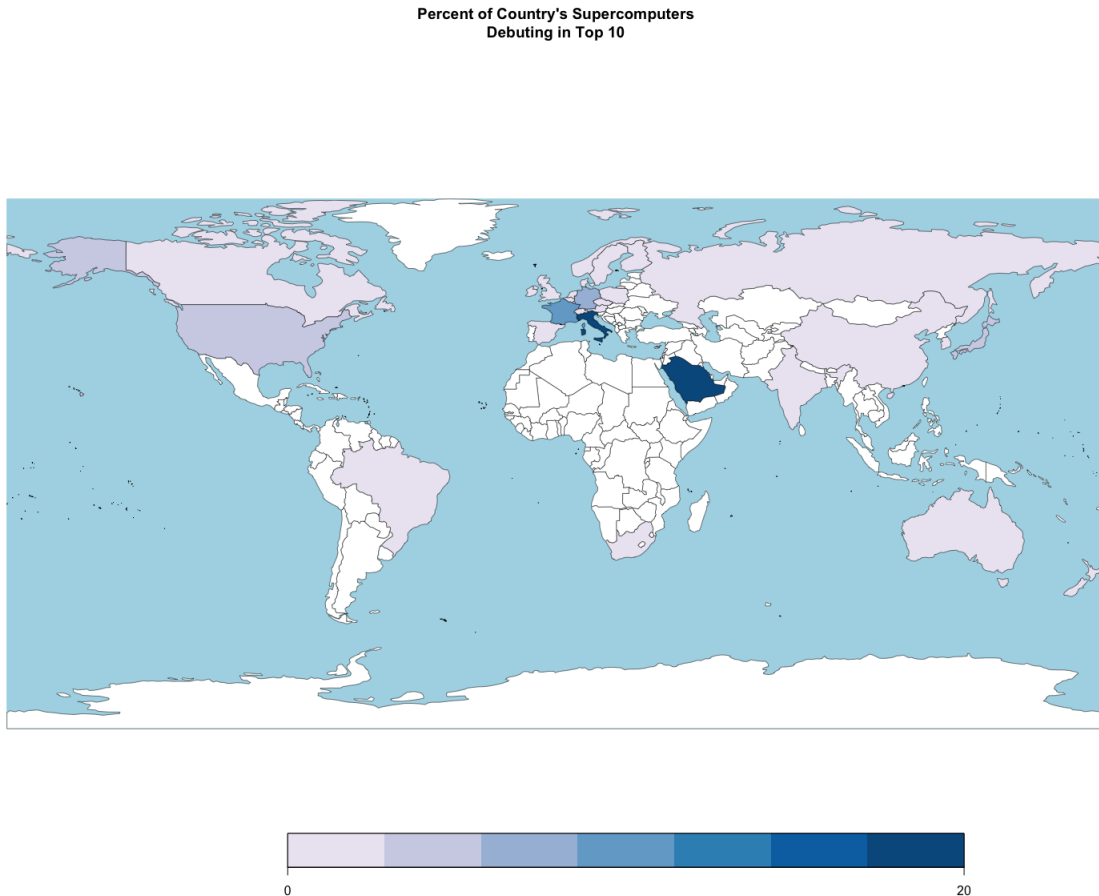


Figure 1: Percent of each country's supercomputers that debuted in the top 10.

The second set of data shows that in terms of applications, researchers and vendors may have a little bit more pride in their supercomputer rankings than those in academia or general industry. Vendors have a major incentive to have the best supercomputer ("look, our chips are the best!"), whereas academia doesn't (they just need it to work, and will take what funding they can get from NSF to get things done). An interesting note here is that there is only one on the list that is classified, and it debuted at rank 150. As for narrative with that one, it would seem that governments doing classified work keep their supercomputers' specifications secret as to not reveal their actual capabilities. So that's fun.

Country	Value
Italy	20.00
Saudi Arabia	20.00
France	11.11
Germany	7.69
United States	5.45
Japan	3.45
China	2.38

Table 1: Percent of each country’s supercomputers that debuted in the top 10. Countries on the top 500 that have not debuted in the top 10: Australia, Austria, Belgium, Brazil, Canada, Czech Republic, Denmark, Finland, India, Ireland, Korea, South, Netherlands, New Zealand, Norway, Poland, Russia, Singapore, South Africa, Spain, Sweden, Switzerland, United Kingdom.

Segment	Value
Research	25.23
Vendor	14.29
Government	11.63
Academic	9.57
Industry	1.23
Classified	0.00

Table 2: Percent of segment’s supercomputers debuting in top 20.




Figure 2: Percent of segment’s supercomputers debuting in top 20.

```

1 library(rworldmap)
2 library(reshape2)
3 library(xtable)
4
5 top500 = read.csv("/Users/Evan/Dropbox/A-Grad-School/CSE380/TOP500_201606.csv")
6 countryData = 100 * colSums(xtabs(~ top500$First.Rank + top500$Country)[1:9,]) /
7   colSums(xtabs(~ top500$First.Rank + top500$Country))
8
9 countryData = melt(countryData)
10
11 df = data.frame(Country = labels(countryData)[[1]],
12               Value = countryData[[1]])
13
14 #join to a coarse resolution map
15 spdf <- joinCountryData2Map(df,
16                             joinCode = "NAME",
17                             nameJoinColumn = "Country")
18
19 mapCountryData(spdf,
20               nameColumnToPlot = "Value",
21               catMethod = "fixedWidth",
22               oceanCol = 'lightblue',
23               colourPalette = brewer.pal(9, "PuBu")[2:8],
24               missingCountryCol = '#ffffff',
25               borderCol = 'black',
26               mapTitle = 'Percent of Country\'s Supercomputers\nDebuting in Top 10')
27
28 print(xtable(subset(df[order(df$Value, decreasing = TRUE),], Value > 0)), include.rownames = FALSE)
29
30
31
32
33 segmentData = 100 * colSums(xtabs(~ top500$First.Rank + top500$Segment)[1:19,]) /
34   colSums(xtabs(~ top500$First.Rank + top500$Segment))
35
36 segmentData = melt(segmentData)
37
38 segmentdf = data.frame(Segment = labels(segmentData)[[1]],
39                       Value = segmentData[[1]])
40 barplot(segmentdf$Value,
41         names.arg = segmentdf$Segment,
42         col = '#7777ee',
43         main = "Percent of Segment\'s Supercomputers\nDebuting in Top 20", las=2)
44
45 print(xtable(segmentdf[order(segmentdf$Value, decreasing = TRUE), ], include.rownames = FALSE)

```

0.1

A Stampede core is “faster” than a Lonestar 4 core because it can do twice as many floating-point operations per clock period. To show a full comparison:

System	Clock Speed (GHz)	flops / cp	Gflops / second (1 core)	cores / node	Gflops / node (all cores)
Lonestar 4	3.33	4	13.32	12	159.84
Lonestar 5	2.6	8	20.8	24	499.2
Stampede	2.7	8	21.6	16	345.6

So in terms of floating-point operations per second (per core), Stampede is faster than Lonestar 5 is faster than Lonestar 4. However, in terms of peak performance per node, Lonestar 5 actually wins, followed by Stampede, then by Lonestar 4. All this even though Lonestar 4 has the fastest clock speed.

0.2

- How many operations per clock cycle?
 - Intel Xeon E5-2680: 8 Flop / cp
- Clock rate?
 - 2.7 GHz
- $8 \text{ Flop / cp} * 2.7 \text{ GHz} = 21.6 \text{ GFlops}$
- 8-way set associative L1 cache [1]. Assume 2 W / cp.
 - $2 * 2.7 * 8 = 43.2 \text{ GB/s}$
- Assume $\sim 0.25 \text{ W/cp}$ from main memory
 - $0.25 * 2.7 * 8 = 5.4 \text{ GB/s}$

0.3

```
# getent passwd | cut -f 7 -d: | sort | uniq
```

The above command works as follows:

- `bash` reads the command, finding no splats or parts of the command that need to be filled in.
- `getent` is called to query the Name Service Switch libraries to get the `passwd` database
- `getent` returns (to standard output) a line for each user including username, group, name, home directory and shell, separated by colons.
- The pipe redirects this standard output to standard input for `cut`, which takes the seventh field as separated by the “:” delimiter, which in the case of `getent passwd` is the location of the shell, and prints it to standard output.
- The next pipe redirects to `sort`, which sorts the list alphanumerically. It outputs the same number of lines as the input, but sorted.
- The final pipe redirects to `uniq` which merges adjacent matching lines. As such, since the lines are already sorted and all the matching lines will be next to each other, `uniq` will output each shell used on stampede exactly once.

0.4

```
# getent passwd | cut -f 7 -d: | sort | uniq --count | sort -bgr
11504 /bin_bash
8949 /bin/bash
136 /bin_tcsh
80 /bin/tcsh
31 /sbin/nologin
22 /bin/csh
20 /bin_csh
18 /bin/zsh
13 /bin_zsh
1 /sbin/shutdown
1 /sbin/halt
1 /bin/sync
```

So 8948 active users on Stampede use bash, 80 use tcsh and the next most popular actual shell is csh.

0.5

```
# getent passwd | cut -f 1 -d: | awk '{print length(), $0}' | sort -g | tail -n1
13 avahi-autoipd
```

avahi-autoipd is the longest TACC username, at 13 characters. It's the Avahi IPv4LL Stack (from `finger avahi-autoipd`), so it's a user, just not necessarily a person.

```
# getent passwd | cut -f 1 -d: | awk '{print length()}' | grep "8" | wc -l
10051
```

10051 users have an 8-character username.

0.6

Question: How many usernames are a person's full name (first and last words in their name)?

```
# getent passwd | cut -f 1,5 -d: --output-delimiter " " | awk '{split($0, a, " "); \
username=a[1]; name = substr(tolower(sprintf("%s%s", a[2], a[length(a)])), 1, 8);\
f1 = index(username, name); f2 = index(name, username); print((f1 == f2) && (f1 == 1))}' \
| grep 1 | wc -l
766
```

0.7

0.7.1

```
# cat words | grep "^[a-zA-Z]" | awk '{print tolower(substr($0, 0, 1))}' | \
uniq --count | sed 's/ //g' | \
awk '{print substr($0, length($0), length($0)) ":" substr($0, 0, length($0) - 1)}'
a:31788
b:25192
c:40042
d:23263
e:17149
f:16263
g:15075
h:19317
i:15377
j:4426
k:6490
l:14680
m:26175
n:16668
o:15274
p:42073
q:3274
r:21688
s:51774
t:25875
u:23930
v:7019
w:12023
x:631
y:1820
z:1914
```

0.7.2

```
# cat words | awk '{print tolower($0)}' | grep "s$" | wc -l
```


There's also \$ARCHIVE which is tape storage through Ranch.

0.9.1

```
# uname -r
2.6.32-431.17.1.el6.x86_64
```

The version is 2.6.32

0.9.2

```
# curl https://www.kernel.org/pub/linux/kernel/v2.6/linux-2.6.32.27.tar.gz > \
linux-2.6.32.27.tar.gz && gunzip linux-2.6.32.27.tar.gz && tar -xvf linux-2.6.32.27.tar
```

(I like seeing all the files as it unarchives them).

0.9.3

```
# du -hc | tail -n1
418M total
# find . | wc -l
32365
# find . | grep '\.c$' | wc -l
13149
```

This version of Linux is 418 megabytes. Including directories (everything is a file or process!), this version has 32365 files, 13149 of which are .c files.

0.9.4

```
# ls -R | grep "^\..*:$" | wc -l
1878
# find . -type d | wc -l
1878
# tree -d -i --noreport | wc -l
1878
```

0.9.5

```
# find . | grep '\.[hc]$$' | xargs wc -l | sort -bg | grep "\.[hc]$" | tail -n1
17722 ./sound/pci/hda/patch_realtek.c
```

The /sound/pci/hda/patch_realtek.c file is the largest (in terms of lines of code) and has 17,722 lines.

0.9.6

```
# find . | grep '\.[hc]$$' | xargs wc -l | tail -n1
828952 total
```

Including comments, there are 828,952 lines of code in this version of the Linux kernel.

0.9.7

Question: How many main functions are there? How many versions have the canonical argument names `int argc` and `char *argv`?

```
# find . | grep '\.[hc]$\ ' | xargs grep "int main(int .*, char .*)"
51
# find . | grep '\.[hc]$\ ' | xargs grep "int main(int argc, char \*argv\[\\])" | wc -l
27
```

References

[1] <http://www.cpu-world.com/CPUs/Xeon/Intel-Xeon%20E5-2680.html>